

Sentiment Analysis on Social Media

CS918
Week 9

Adam Tsakalidis

Module Organiser: Dr Arkaitz Zubiaga



Definition

- **Definition**
 - Given a piece of text, classify it with respect to its sentiment
- **Different levels**
 - Sentence-, phrase-, entity-level
- **Three classes**
 - Positive, negative, neutral
- **Two main approaches**
 - Supervised VS Unsupervised
- **Similar to other classification/clustering tasks**



Background

- Closely related tasks:
 - Emotion Classification
 - Sarcasm Detection
 - Stance Detection
- Applications of SA:
 - Monitoring public mood
 - Election prediction
 - Stock market
 - Many others!

SA Approaches on OSM

- **All-vs-All:**
 - Three classes, one model
 - Most popular approach
 - **One-vs-All**
 - Train on true/false examples for every class
 - One model per class
 - **Two-Step process**
 - Subjectivity detection (neutral VS subjective)
 - Polarity detection (positive VS negative)

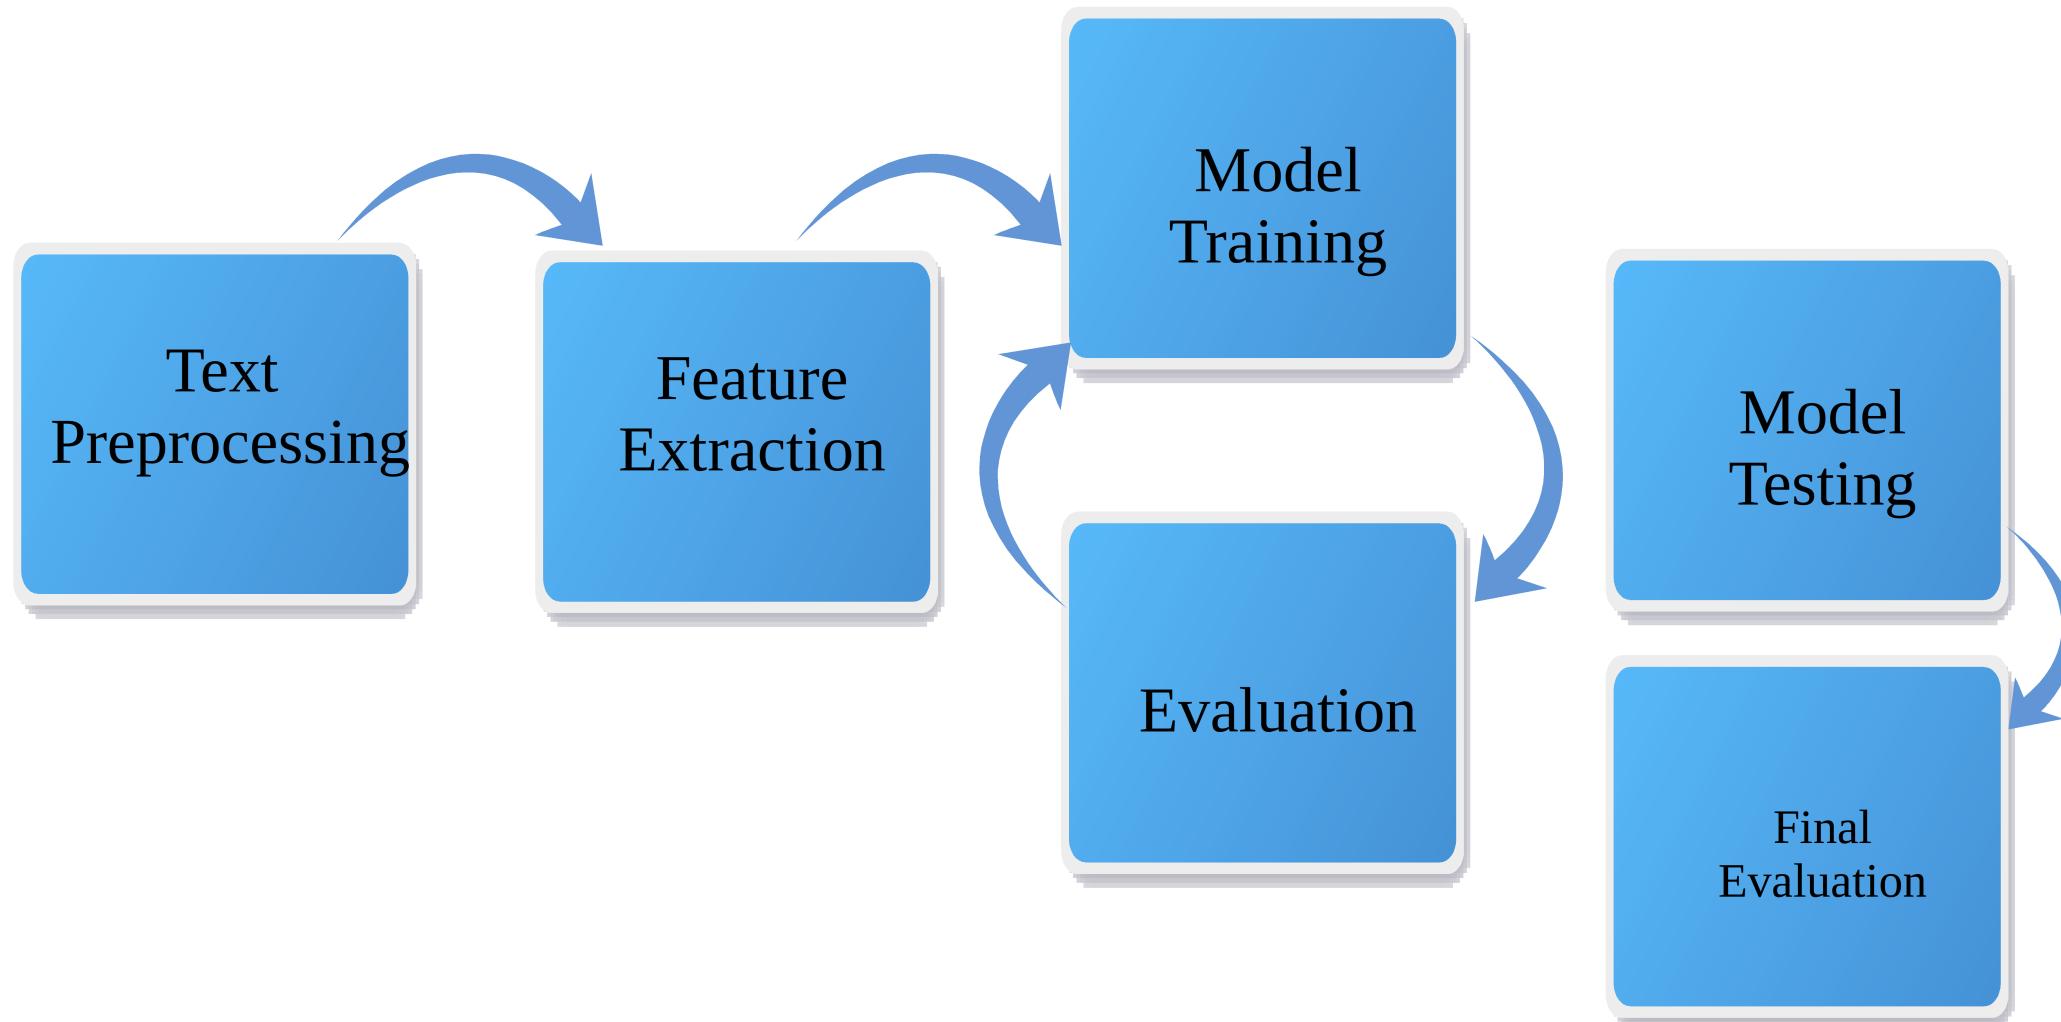
Approaches on SA

- **Supervised Methods**
 - Learn from labelled instances
 - Most popular/successful on OSM
- **Unsupervised Methods**
 - No need for manual annotation
 - Lower performance
- **Semi-Supervised Methods**
 - Manual annotation is expensive
 - Idea: pseudo-label training examples
 - e.g., :) vs :(

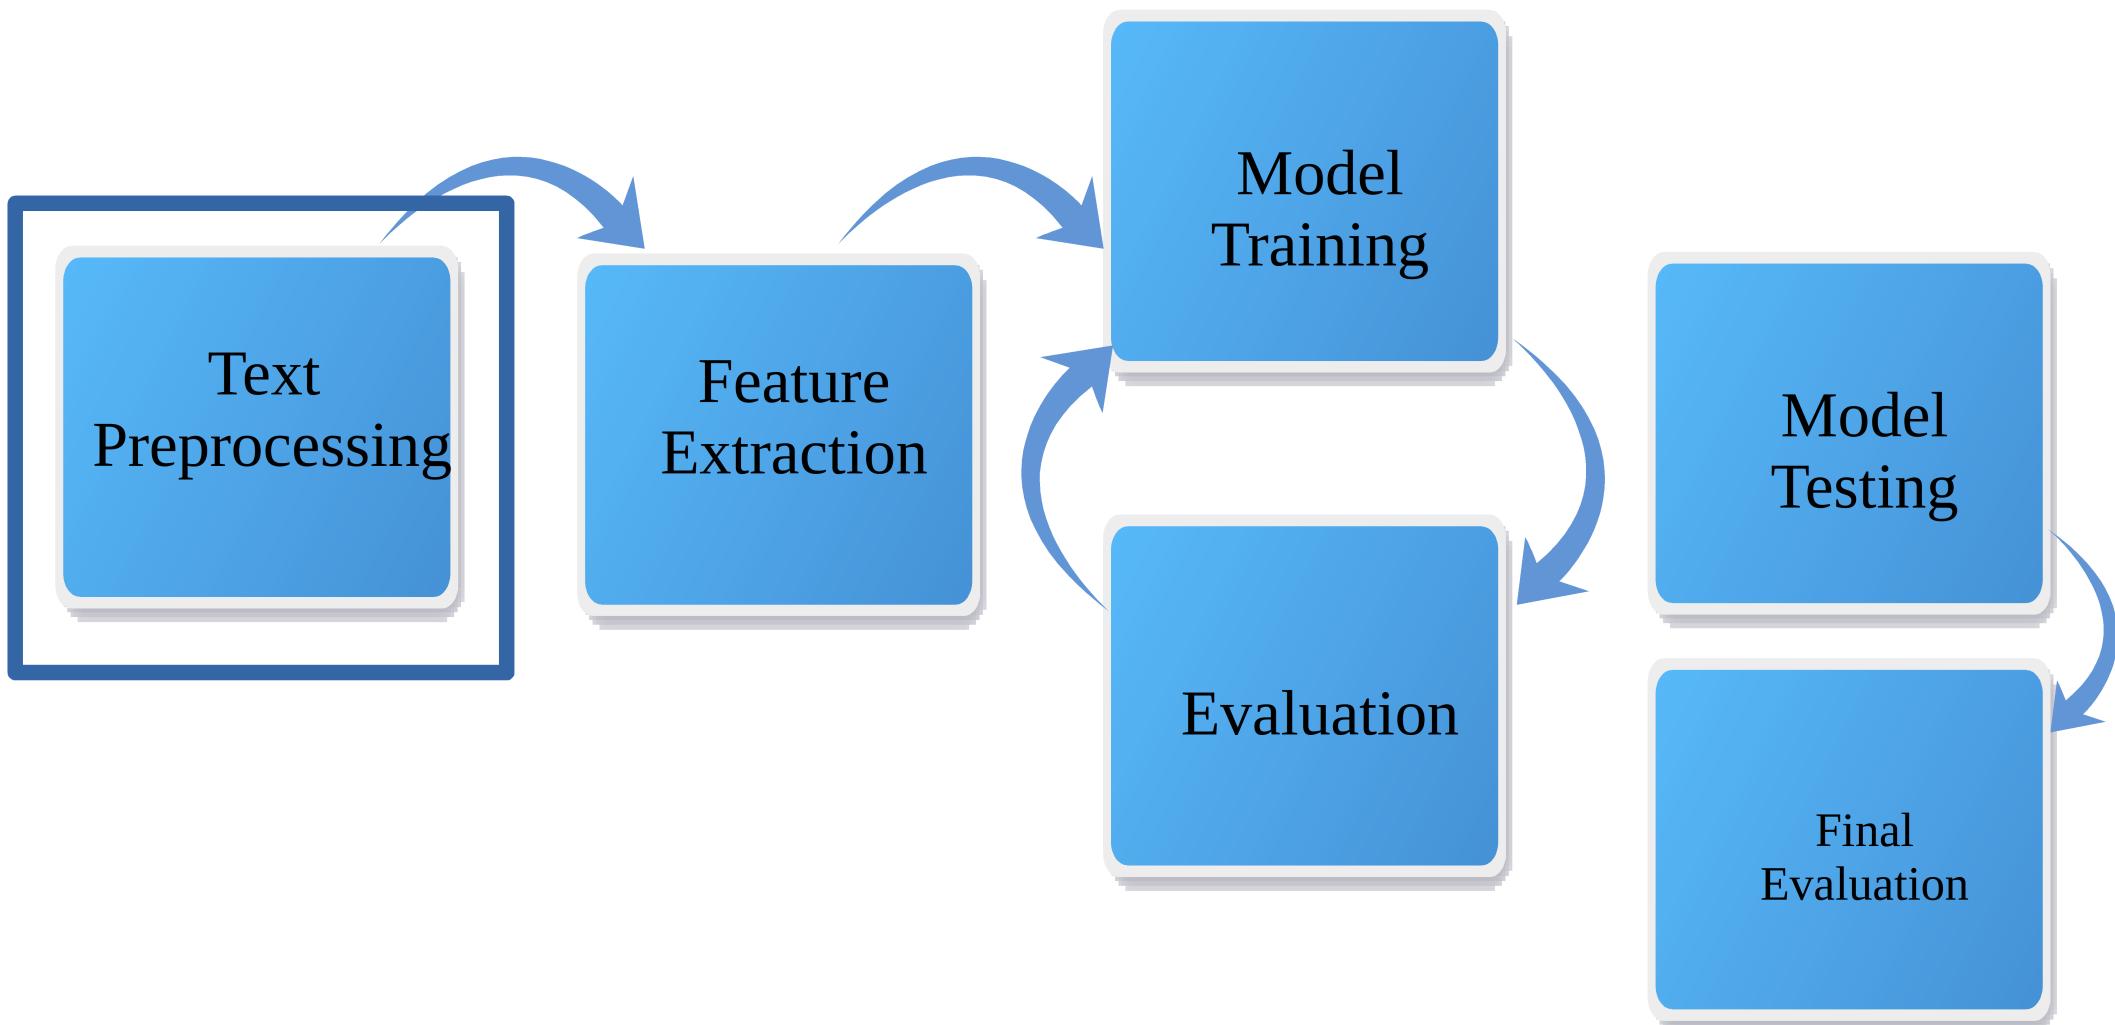
Exercise 2

- **Tweet Sentiment Classification**
 - **Task:** Given a tweet, classify whether the tweet is of positive, negative, or neutral sentiment.
 - **Data:** one training set; one development set; 3 test sets
 - **Evaluation:** macro-average f-score (pos/neg)
 - **Five steps:**
 - **Preprocess** the tweets (re-use your code!)
 - Extract **features** (see next)
 - Train **classifiers** on the extracted features of the training set
 - **Evaluate** using 10-fold cv (train set) or dev set
 - **Apply** best model to the (3) test sets & **evaluate** its performance

Exercise 2



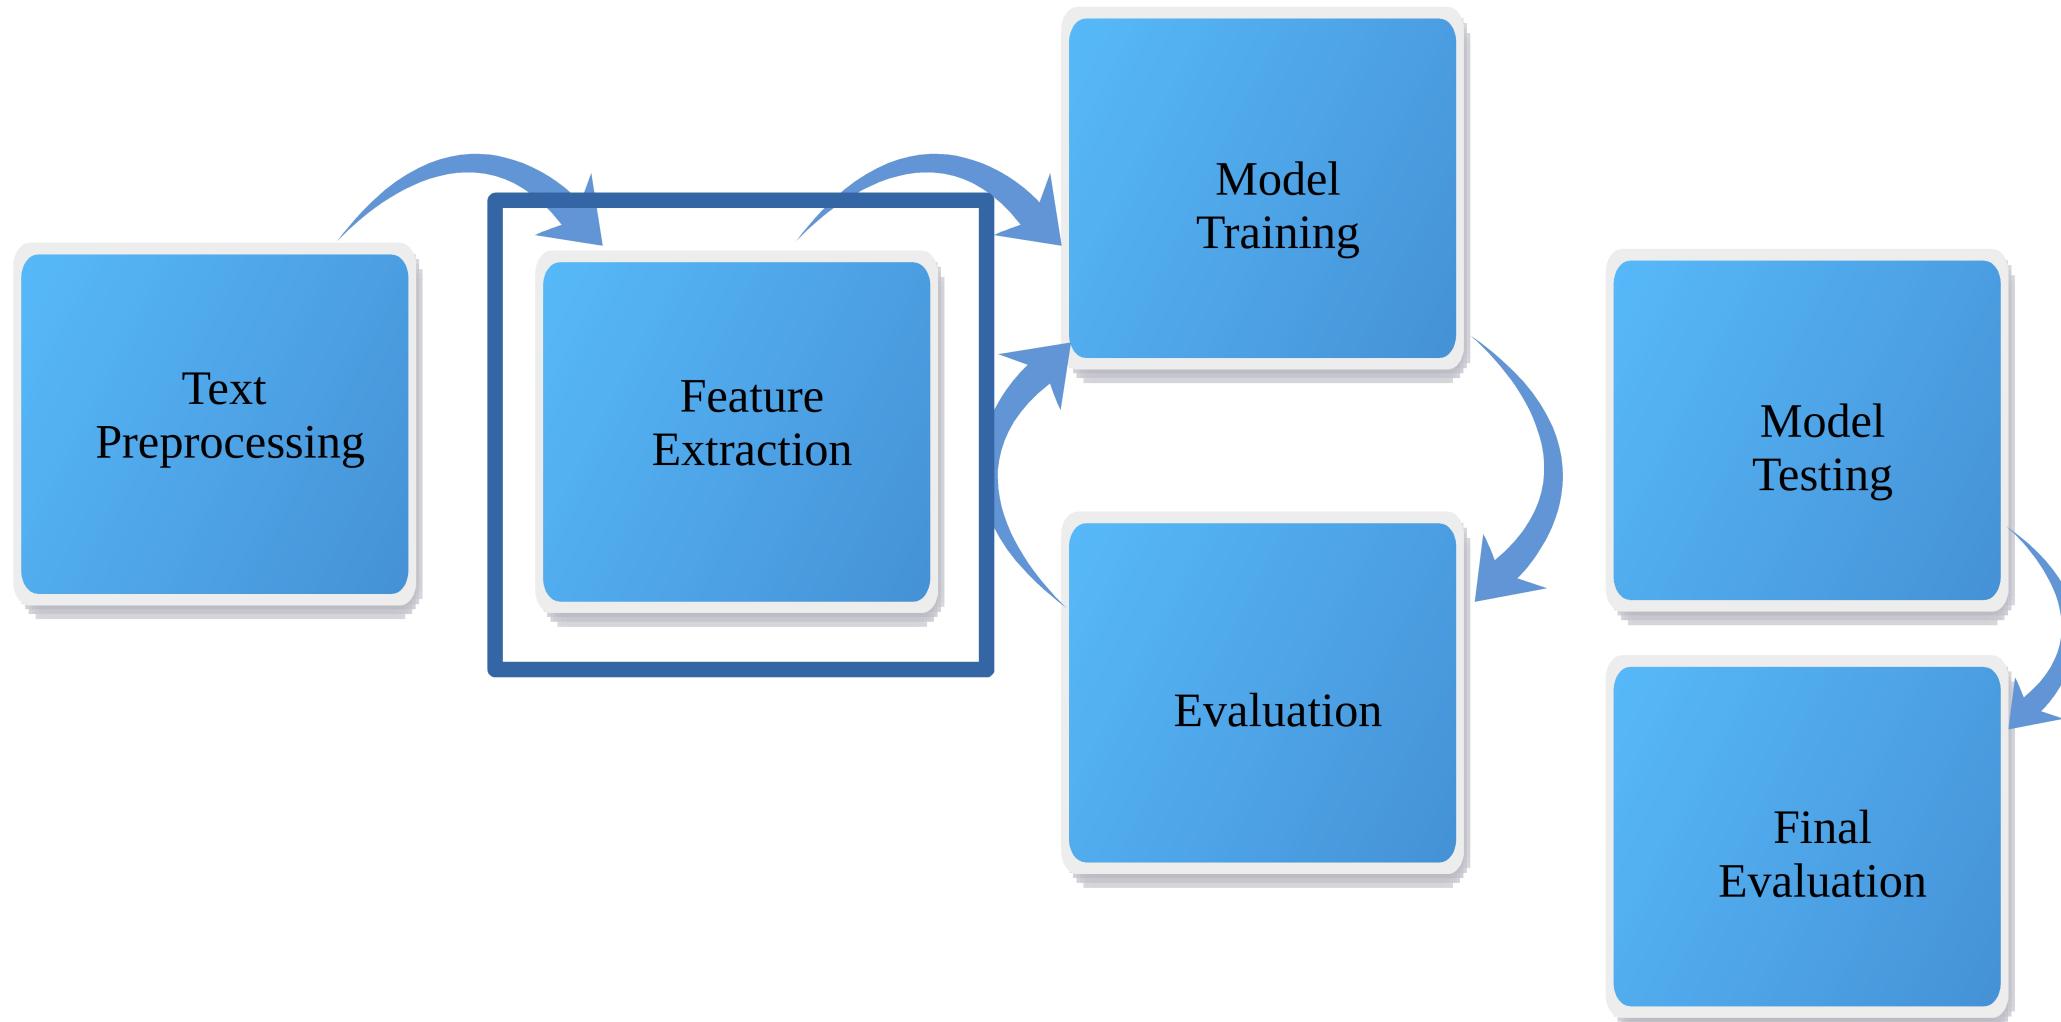
Exercise 2



Step 1: Text preprocessing

- Recap (see previous seminar):
 - **Lowercase** (all-upper-case tokens might be important):
 - “This is NOT good” → “this is uppercasenot good”
 - **Replace elongated words:**
 - “This is so coooooooolllll” → “This is so coooll”
 - **Replace mentions/urls/etc. with identifiers:**
 - “Listening to the new song by @user <https://youtube.com>”
 - “listening to the new song by usrmnt urlink”
 - **Remove non-alphanumeric**
 - **Others:** smileys replacement, stemming

Exercise 2



Step 2: Feature extraction

- Ngrams
 - {binary, tf, tfidf}, {POS}
 - Sentiment Lexicons
 - Counts/summation of scores derived from lexicons
 - Word Embeddings
 - Use pre-trained vectors
 - Use functions over vector dimensions
 - Twitter-specific
 - Hashtags, uppercase words, tweet length...

Step 2: ngrams

- Doc#1 (train): “great piece of art have a look”
- Doc#2 (train): “omg this notdoes look great”
- Doc#3 (test): “omg is this real”
- Vocabulary (e.g., binary unigrams) based on the **training set**:

| Doc | a | art | great | have | look | notdoes | of | omg | piece | this |
|-------|---|-----|-------|------|------|---------|----|-----|-------|------|
| Doc#1 | 1 | 1 | 1 | 1 | 1 | 0 | 1 | 0 | 1 | 0 |
| Doc#2 | 0 | 0 | 1 | 0 | 1 | 1 | 0 | 1 | 0 | 1 |
| Doc#3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 |

- POS tagging: concatenate word+POS when building the vocabulary (e.g., “greatADJ”)
- Resources: sklearn (see the vectorisers: <http://bit.ly/2FPfiCt>)

Step 2: ngrams

- Doc#1 (train): “great piece of art have a look”
- Doc#2 (train): “omg this notdoes look great”
- Doc#3 (test): “omg **is** this **real**” ←
- Vocabulary (e.g., binary unigrams) based on the **training set**:

| Doc | a | art | great | have | look | notdoes | of | omg | piece | this |
|-------|---|-----|-------|------|------|---------|----|-----|-------|------|
| Doc#1 | 1 | 1 | 1 | 1 | 1 | 0 | 1 | 0 | 1 | 0 |
| Doc#2 | 0 | 0 | 1 | 0 | 1 | 1 | 0 | 1 | 0 | 1 |
| Doc#3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 |

- POS tagging: concatenate word+POS when building the vocabulary (e.g., “greatADJ”)
- Resources: sklearn (see the vectorisers: <http://bit.ly/2FPfiCt>)

Step 2: word embeddigs

- **Inputs:**
 - preprocessed tweet with n words;
 - word embeddings: {word→[vals]}
 - **Output:** tweet with n -by- D scores (D : dimension of embeddings)
 - **Features:** functions over each of the D resulting dimensions
 - average/max/min/multiplication/summation
 - each applied function generates D features for a tweet
 - **Resources** (pre-trained word vectors):
 - Glove: <https://nlp.stanford.edu/projects/glove/>
 - word2vec: <https://code.google.com/archive/p/word2vec/>
 - Others: <http://ir.hit.edu.cn/~dytang/>

Step 2: sentiment lexicons

- **Inputs:**
 - preprocessed tweet with n words;
 - lexicon_1: {word→[vals]} (e.g., D-dimensional)
 - lexicon_2: {word→label} (e.g., binary)
 - ...
 - **Features:**
 - D-dimensional (lexicon_1): summation of each dimension
 - 2-dimensional (lexicon_2): number of (pos/neg) words
 - ...
 - Use the concatenation of all of these features

Step 2: sentiment lexicons

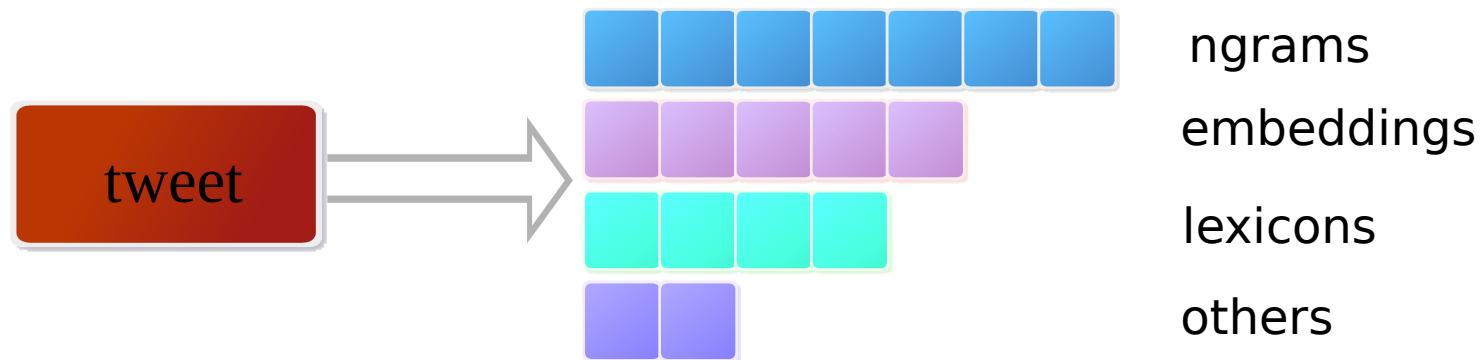
- Resources:
 - <http://sentiwordnet.isti.cnr.it/>
 - <https://www.cs.uic.edu/~liub/FBS/sentiment-analysis.html>
 - http://mpqa.cs.pitt.edu/lexicons/subj_lexicon/
 - <http://www.wjh.harvard.edu/~inquirer/>
 - <http://liwc.wpengine.com/>
 - Beware: different feature type per lexicon!
 - Prefer counts for binary lexicons
 - Prefer overall (or per-class) sum for lexicons with real values

Step 2: other features

- Twitter-specific features:
 - #elongated words
 - #all-upper-case words
 - presence/absence of negation
 - presence/absence of URL (binary)
 - presence/absence of user mention (binary)
 - length of tweet (in chars or words)
 - #POS tags (e.g., #adjectives)
 - ...
- Use the concatenation of all of these features

Step 2: Summary

- Extract different feature sets



- Use them individually or in concatenation
 - e.g., ngrams+embeddings+lexicons, lexicons+embeddings...
- Feed them into a classifier (see next)
- Validate on train/dev set; select best features/classifier
- If your features are bad, your model will be bad!

Step 2: example

"I think this is AWESOME! #music <http://url.com>"

"@user123 This is actually soooo bad..."

| word | label |
|---------|-------|
| good | pos |
| bad | neg |
| awesome | pos |
| awful | neg |

Lexicons

Word
Vectors

| word | dim1 | dim2 | dim3 |
|----------|------|------|------|
| actually | 0.1 | 0.2 | 0.1 |
| awesome | 0.3 | 0.2 | 0.3 |
| bad | 0.8 | 0.6 | 0.6 |
| I | 0.1 | 0.2 | 0.3 |
| is | 0.1 | 0.1 | 0.1 |
| music | 0.3 | 0.8 | 0.2 |
| so | 0.1 | 0.2 | 0.5 |
| think | 0.6 | 0.5 | 0.5 |
| this | 0.1 | 0.2 | 0.8 |

| word | happy | sad | surprise | anger |
|---------|-------|-----|----------|-------|
| soo | 0.3 | 0.1 | 0.2 | 0.2 |
| bad | 0 | 0.9 | 0.2 | 0.4 |
| awesome | 0.9 | 0 | 0.3 | 0 |

Step 2: example

“I think this is AWESOME! #music <http://url.com>”

“@user123 This is actually soooo bad...”

ngram features (assume binary unigrams/bigrams)?

Step 2: example

"I think this is AWESOME! #music <http://url.com>"

"@user123 This is actually soooo bad..."

Doc1 unigrams:

Doc1 bigrams:

Step 2: example

"I think this is AWESOME! #music <http://url.com>"

"@user123 This is actually soooo bad..."

Doc1 unigrams: {i, think, this, is, awesome, music, urlink}

Doc1 bigrams: {i think, think this, this is, is awesome, awesome music,
music urlink}

Step 2: example

"I think this is AWESOME! #music <http://url.com>"

"@user123 This is actually soooo bad..."

Doc1 unigrams: {i, think, this, is, awesome, music, urlink}

Doc1 bigrams: {i think, think this, this is, is awesome, awesome music,
music urlink}

$$|\mathbf{V}|=13$$

| | i | think | this | is | awesome | music | urlink | i think | think this | this is | is awesome | awesome music |
|----|---|-------|------|----|---------|-------|--------|---------|------------|---------|------------|---------------|
| #1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| #2 | | | | | | | | | | | | |

Step 2: example

"I think this is AWESOME! #music <http://url.com>"

"@user123 This is actually soooo bad..."

Doc1 unigrams: {i, think, this, is, awesome, music, urlink}

Doc1 bigrams: {i think, think this, this is, is awesome, awesome music, music urlink}

$$|\mathbf{V}|=13$$

| | i | think | this | is | awesome | music | urlink | i think | think this | this is | is awesome | awesome music |
|----|---|-------|------|----|---------|-------|--------|---------|------------|---------|------------|---------------|
| #1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| #2 | | | | 1 | 1 | | | | | 1 | | |

Step 2: example

"I think this is AWESOME! #music <http://url.com>"

"@user123 This is actually soooo bad..."

| word | label |
|---------|-------|
| good | pos |
| bad | neg |
| awesome | pos |
| awful | neg |

| word | happy | sad | surprise | anger |
|---------|-------|-----|----------|-------|
| soo | 0.3 | 0.1 | 0.2 | 0.2 |
| bad | 0 | 0.9 | 0.2 | 0.4 |
| awesome | 0.9 | 0 | 0.3 | 0 |

Which lexicon features would you use?

Step 2: example

"I think this is AWESOME! #music <http://url.com>"

"@user123 This is actually soooo bad..."

| word | label |
|---------|-------|
| good | pos |
| bad | neg |
| awesome | pos |
| awful | neg |

| word | happy | sad | surprise | anger |
|---------|-------|-----|----------|-------|
| soo | 0.3 | 0.1 | 0.2 | 0.2 |
| bad | 0 | 0.9 | 0.2 | 0.4 |
| awesome | 0.9 | 0 | 0.3 | 0 |

| | | | | | | | |
|----|---|---|---|---|---|---|---|
| | ? | ? | ? | ? | ? | ? | ? |
| #1 | ? | ? | ? | ? | ? | ? | ? |
| #2 | ? | ? | ? | ? | ? | ? | ? |

Step 2: example

“I think this is AWESOME! #music <http://url.com>”

“@user123 This is actually soooo bad...”

| word | label |
|---------|-------|
| good | pos |
| bad | neg |
| awesome | pos |
| awful | neg |

| word | happy | sad | surprise | anger |
|---------|-------|-----|----------|-------|
| soo | 0.3 | 0.1 | 0.2 | 0.2 |
| bad | 0 | 0.9 | 0.2 | 0.4 |
| awesome | 0.9 | 0 | 0.3 | 0 |

| | pos | neg | happy | sad | surprise | anger |
|----|-----|-----|-------|-----|----------|-------|
| #1 | | | | | | |
| #2 | | | | | | |

Step 2: example

"I think this is AWESOME! #music <http://url.com>"

"@user123 This is actually soooo bad..."

| word | label |
|---------|-------|
| good | pos |
| bad | neg |
| awesome | pos |
| awful | neg |

| word | happy | sad | surprise | anger |
|---------|-------|-----|----------|-------|
| soo | 0.3 | 0.1 | 0.2 | 0.2 |
| bad | 0 | 0.9 | 0.2 | 0.4 |
| awesome | 0.9 | 0 | 0.3 | 0 |

| | pos | neg | happy | sad | surprise | anger |
|----|-----|-----|-------|-----|----------|-------|
| #1 | 1 | 0 | 0.9 | 0.0 | 0.3 | 0.0 |
| #2 | 0 | 1 | 0.3 | 1.0 | 0.4 | 0.6 |

Step 2: example

“I think this is AWESOME! #music <http://url.com>”

“@user123 This is actually soooo bad...”

| word | dim1 | dim2 | dim3 |
|----------|------|------|------|
| actually | 0.1 | 0.2 | 0.1 |
| awesome | 0.3 | 0.2 | 0.3 |
| bad | 0.8 | 0.6 | 0.6 |
| I | 0.1 | 0.2 | 0.3 |
| is | 0.1 | 0.1 | 0.1 |
| music | 0.3 | 0.8 | 0.2 |
| so | 0.1 | 0.2 | 0.5 |
| think | 0.6 | 0.5 | 0.5 |
| this | 0.1 | 0.2 | 0.8 |

Let's extract some word embeddings features...

(Remember: functions over dimensions)

Step 2: example

“I think this is AWESOME! #music <http://url.com>”

“@user123 This is actually soooo bad...”

| word | dim1 | dim2 | dim3 |
|----------|------|------|------|
| actually | 0.1 | 0.2 | 0.1 |
| awesome | 0.3 | 0.2 | 0.3 |
| bad | 0.8 | 0.6 | 0.6 |
| I | 0.1 | 0.2 | 0.3 |
| is | 0.1 | 0.1 | 0.1 |
| music | 0.3 | 0.8 | 0.2 |
| so | 0.1 | 0.2 | 0.5 |
| think | 0.6 | 0.5 | 0.5 |
| this | 0.1 | 0.2 | 0.8 |

| Doc#1 | Dim1 | Dim2 | Dim3 |
|---------|------|------|------|
| I | 0.1 | 0.2 | 0.3 |
| think | 0.6 | 0.5 | 0.5 |
| this | 0.1 | 0.2 | 0.8 |
| is | 0.1 | 0.1 | 0.1 |
| awesome | 0.3 | 0.2 | 0.3 |
| music | 0.3 | 0.8 | 0.2 |
| urlink | — | — | — |

Step 2: example

“I think this is AWESOME! #music <http://url.com>”

“@user123 This is actually soooo bad...”

| word | dim1 | dim2 | dim3 |
|----------|------|------|------|
| actually | 0.1 | 0.2 | 0.1 |
| awesome | 0.3 | 0.2 | 0.3 |
| bad | 0.8 | 0.6 | 0.6 |
| I | 0.1 | 0.2 | 0.3 |
| is | 0.1 | 0.1 | 0.1 |
| music | 0.3 | 0.8 | 0.2 |
| so | 0.1 | 0.2 | 0.5 |
| think | 0.6 | 0.5 | 0.5 |
| this | 0.1 | 0.2 | 0.8 |

| Doc#1 | Dim1 | Dim2 | Dim3 |
|------------|------|------|------|
| I | 0.1 | 0.2 | 0.3 |
| think | 0.6 | 0.5 | 0.5 |
| this | 0.1 | 0.2 | 0.8 |
| is | 0.1 | 0.1 | 0.1 |
| awesome | 0.3 | 0.2 | 0.3 |
| music | 0.3 | 0.8 | 0.2 |
| urlink | — | — | — |
| MIN | | | |
| SUM | | | |

Step 2: example

"I think this is AWESOME! #music <http://url.com>"

"@user123 This is actually soooo bad..."

| word | dim1 | dim2 | dim3 |
|----------|------|------|------|
| actually | 0.1 | 0.2 | 0.1 |
| awesome | 0.3 | 0.2 | 0.3 |
| bad | 0.8 | 0.6 | 0.6 |
| I | 0.1 | 0.2 | 0.3 |
| is | 0.1 | 0.1 | 0.1 |
| music | 0.3 | 0.8 | 0.2 |
| so | 0.1 | 0.2 | 0.5 |
| think | 0.6 | 0.5 | 0.5 |
| this | 0.1 | 0.2 | 0.8 |

| Doc#1 | Dim1 | Dim2 | Dim3 |
|------------|------------|------------|------------|
| I | 0.1 | 0.2 | 0.3 |
| think | 0.6 | 0.5 | 0.5 |
| this | 0.1 | 0.2 | 0.8 |
| is | 0.1 | 0.1 | 0.1 |
| awesome | 0.3 | 0.2 | 0.3 |
| music | 0.3 | 0.8 | 0.2 |
| urlink | — | — | — |
| MIN | 0.1 | 0.1 | 0.1 |
| SUM | 1.5 | 2.0 | 2.2 |

Step 2: example

"I think this is AWESOME! #music <http://url.com>"

"@user123 This is actually soooo bad..."

| word | dim1 | dim2 | dim3 |
|----------|------|------|------|
| actually | 0.1 | 0.2 | 0.1 |
| awesome | 0.3 | 0.2 | 0.3 |
| bad | 0.8 | 0.6 | 0.6 |
| I | 0.1 | 0.2 | 0.3 |
| is | 0.1 | 0.1 | 0.1 |
| music | 0.3 | 0.8 | 0.2 |
| so | 0.1 | 0.2 | 0.5 |
| think | 0.6 | 0.5 | 0.5 |
| this | 0.1 | 0.2 | 0.8 |

| #1 | 0.1 | 0.1 | 0.1 | 1.5 | 2.0 | 2.2 |
|----|-----|-----|-----|-----|-----|-----|
| #2 | 0.1 | 0.1 | 0.1 | 1.1 | 1.1 | 1.6 |

Step 2: example

"I think this is AWESOME! #music <http://url.com>"

"@user123 This is actually soooo bad..."

| word | dim1 | dim2 | dim3 |
|----------|------|------|------|
| actually | 0.1 | 0.2 | 0.1 |
| awesome | 0.3 | 0.2 | 0.3 |
| bad | 0.8 | 0.6 | 0.6 |
| I | 0.1 | 0.2 | 0.3 |
| is | 0.1 | 0.1 | 0.1 |
| music | 0.3 | 0.8 | 0.2 |
| so | 0.1 | 0.2 | 0.5 |
| think | 0.6 | 0.5 | 0.5 |
| this | 0.1 | 0.2 | 0.8 |

| #1 | 0.1 | 0.1 | 0.1 | 1.5 | 2.0 | 2.2 |
|----|-----|-----|-----|-----|-----|-----|
| #2 | 0.1 | 0.1 | 0.1 | 1.1 | 1.1 | 1.6 |

Any other features?

Step 2: example

“I think this is AWESOME! #music <http://url.com>”

“@user123 This is actually soooo bad...”

| | length | link | mention | exclmtn | negation | smiley | frown | elong |
|----|--------|------|---------|---------|----------|--------|-------|-------|
| #1 | 7 | 1 | 0 | 1 | 0 | 0 | 0 | 0 |
| #2 | 6 | 0 | 1 | 0 | 0 | 0 | 0 | 1 |

Step 2: example

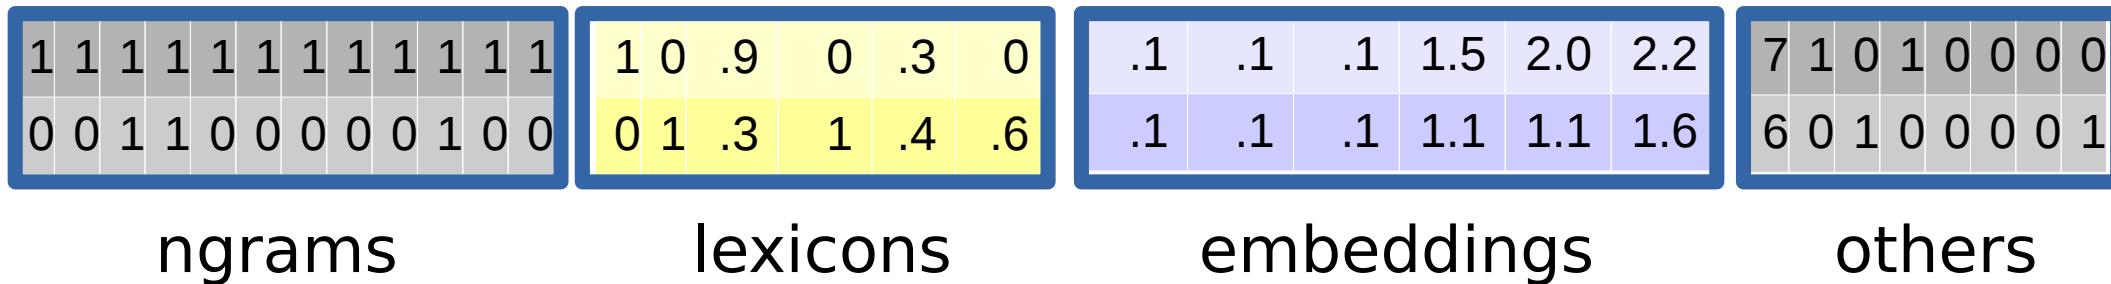
“I think this is AWESOME! #music <http://url.com>”

“@user123 This is actually soooo bad...”

Step 2: example

"I think this is AWESOME! #music <http://url.com>"

"@user123 This is actually soooo bad..."



Use these feature sources individually

Step 2: example

"I think this is AWESOME! #music <http://url.com>"

“@user123 This is actually soooo bad...”

ngrams + lexicons

Use these feature sources individually
or concatenated

Step 2: example

"I think this is AWESOME! #music <http://url.com>"

“@user123 This is actually soooo bad...”

lexicons + embeddings + others

Use these feature sources individually
or concatenated

Step 2: example

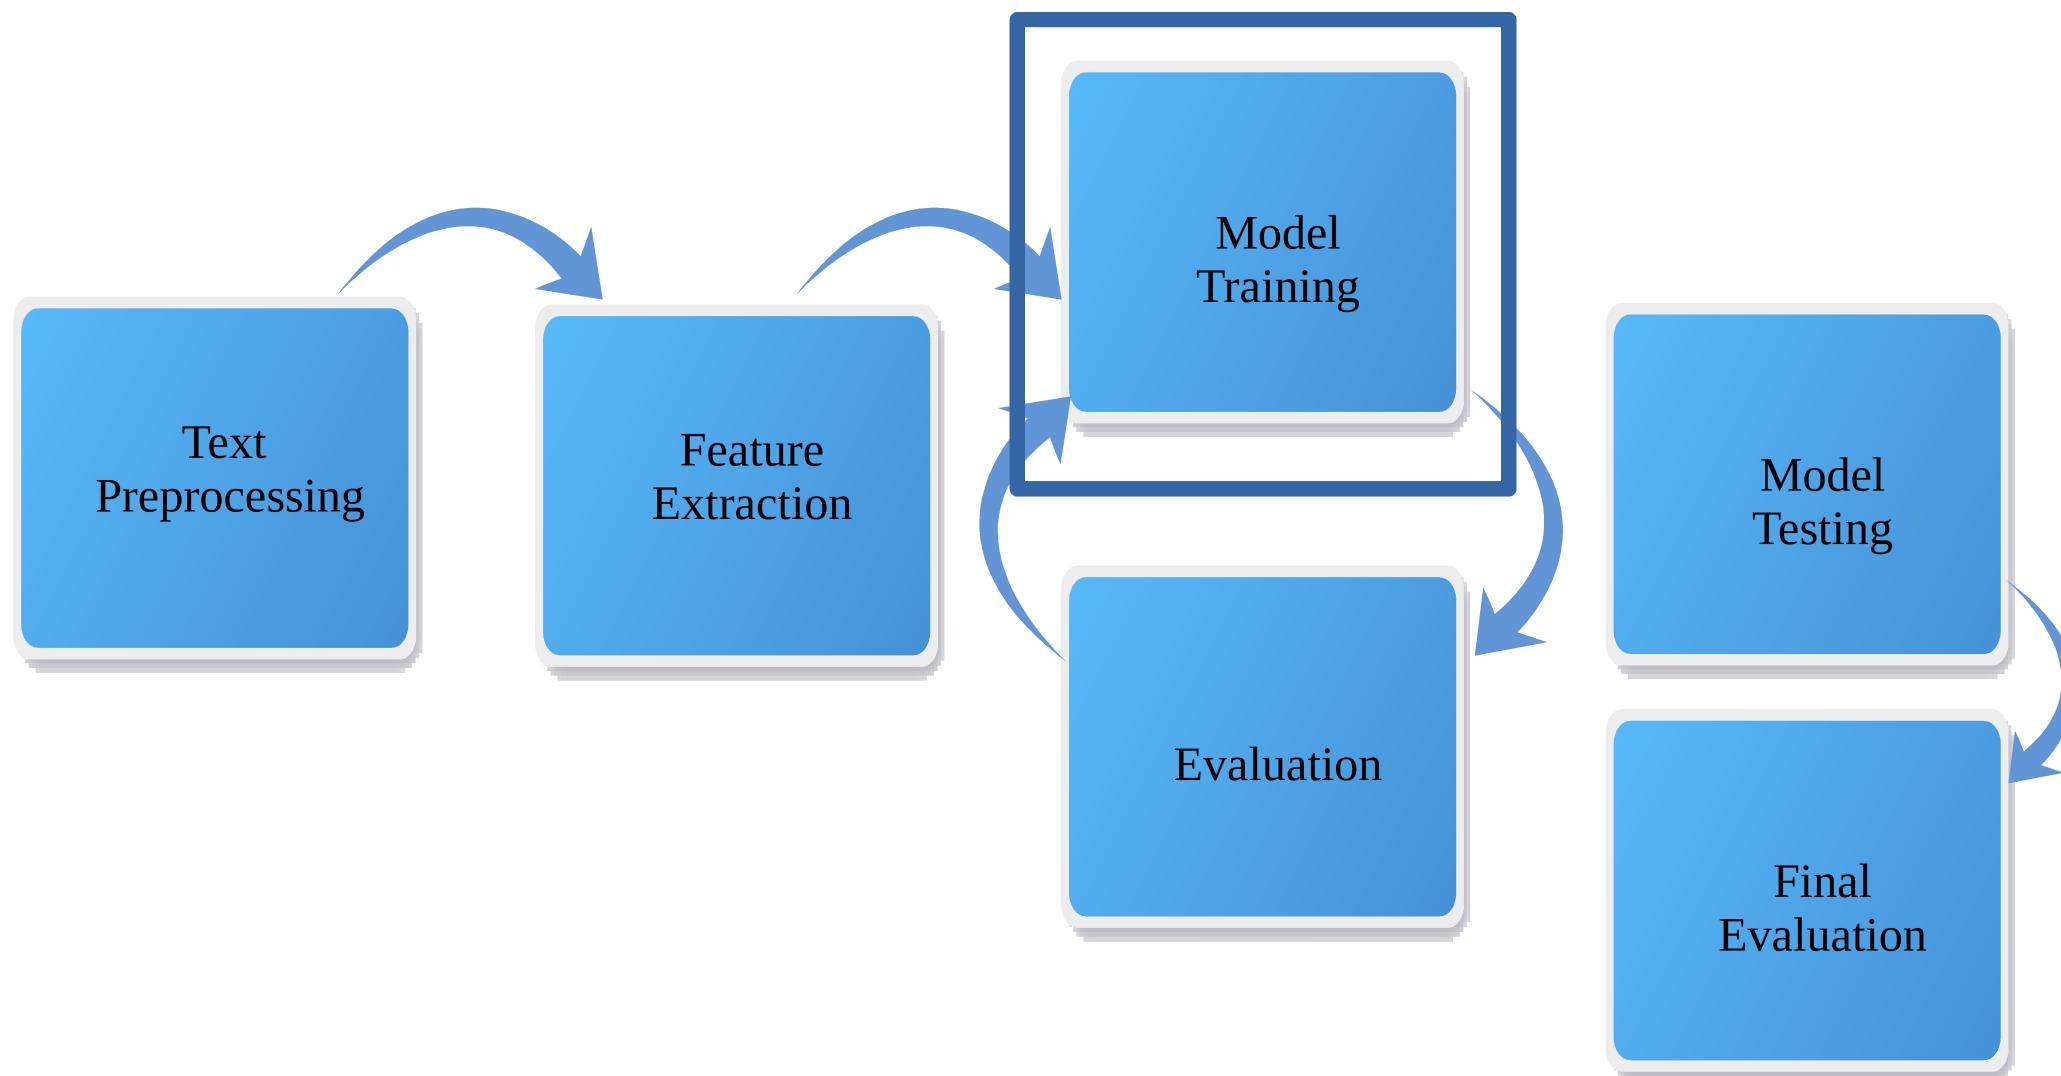
"I think this is AWESOME! #music <http://url.com>"

“@user123 This is actually soooo bad...”

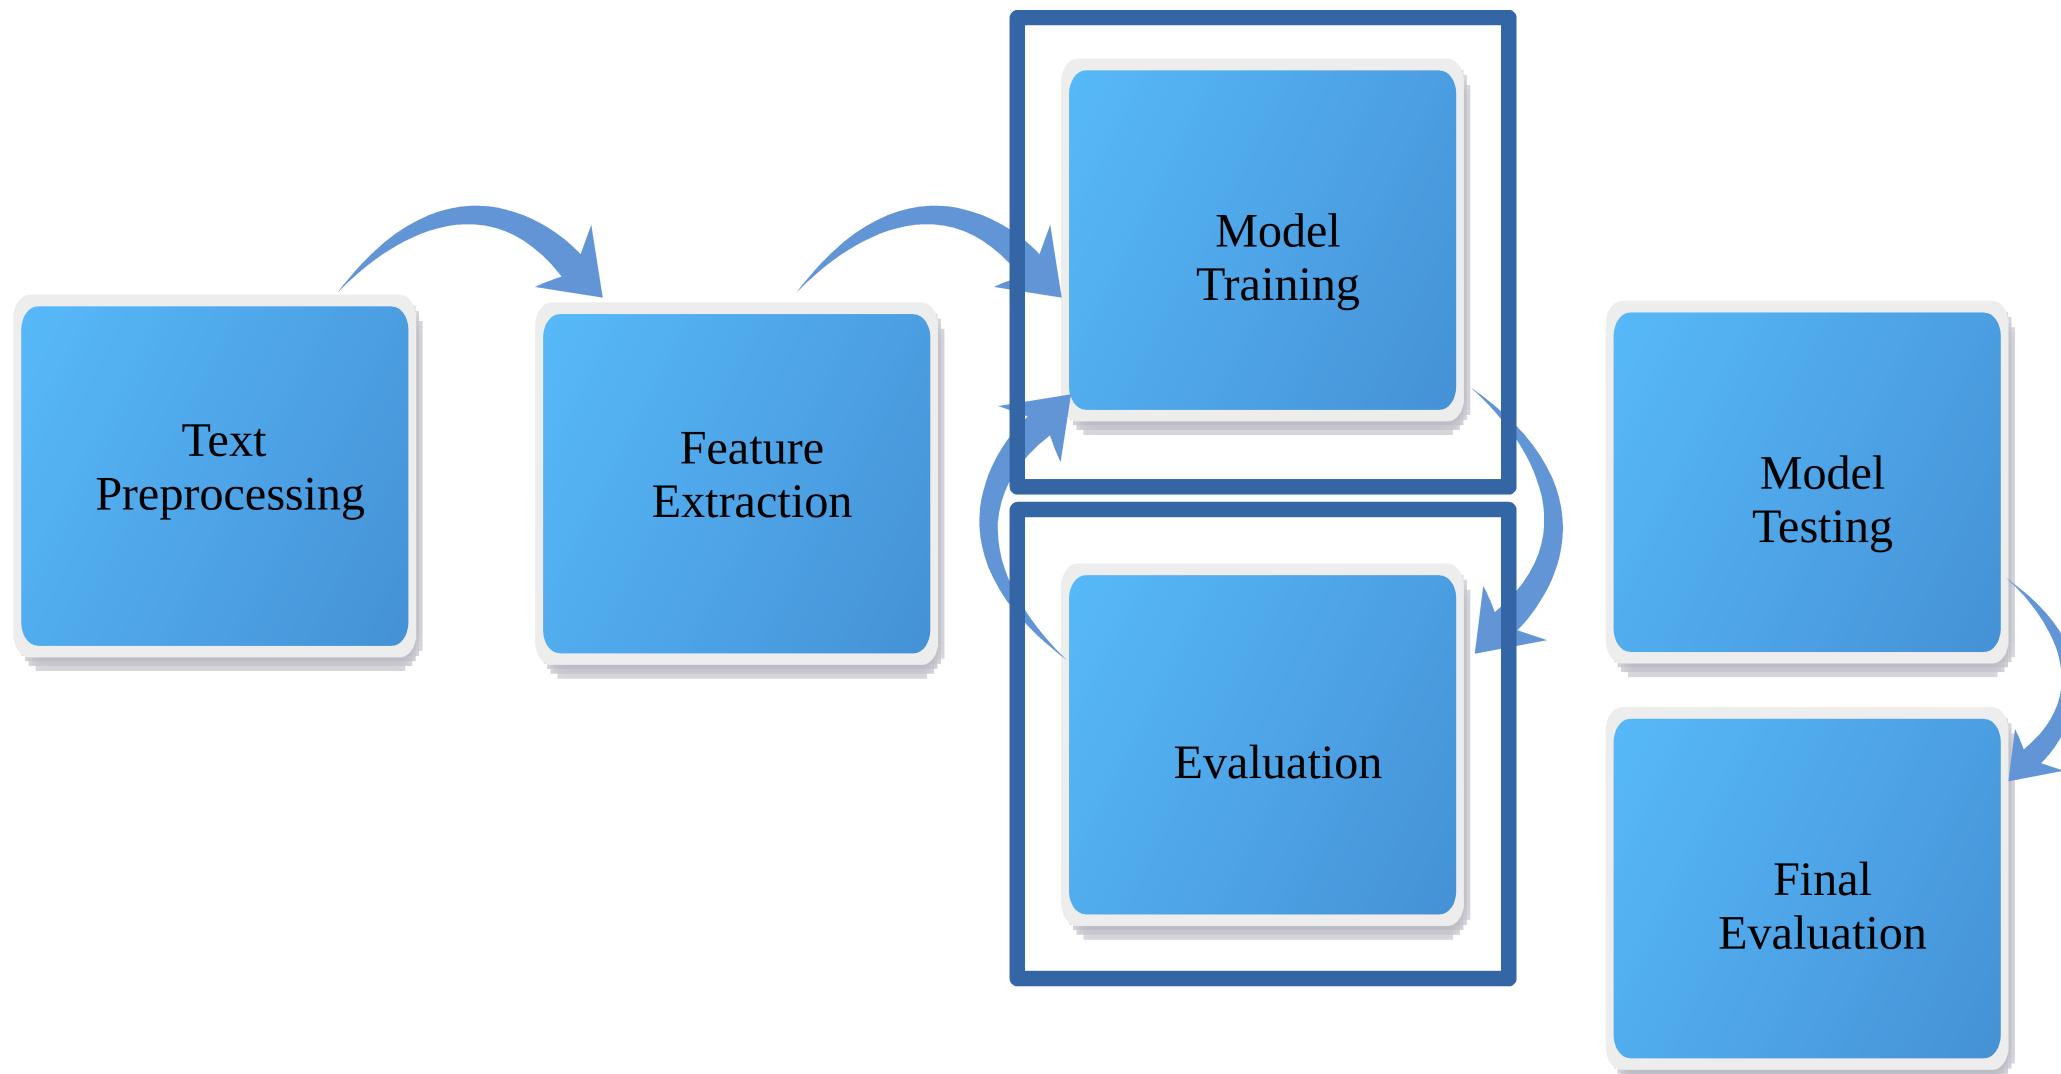
ngrams + lexicons + embeddings + others

Use these feature sources individually
or concatenated

Exercise 2



Exercise 2



Step 3: training

- **Start with simple models**
 - Support Vector Machine
 - Random Forest
 - Logistic Regression
 - Naive Bayes
- **Recommended:**
 - sklearn (<http://scikit-learn.org/stable/>)
- Proceed with deep learning - if time allows to do so

Step 3: training

- Train different classifiers on different feature sets
- Report macro-average f-score (pos/neg), e.g.:

| | ngram | embed | lex | ngram + embed | ngram + lex | embed + lex | ngram + other | lex + other | all |
|----------|-------|-------|-----|---------------|-------------|-------------|---------------|-------------|-----|
| Majority | | | | | | | | | |
| SVM | | | | | | | | | |
| RF | | | | | | | | | |
| LR | | | | | | | | | |

- Questions:
 - How do you select the best parameters per model?
 - How do you select the best model?

Step 3: tuning

- Train all models on the training set
 - Tune parameters: 10-fold cross-validation
 - Regularisers, kernel selection for SVM, number of trees in RF...
 - Hint: GridSearchCV in sklearn
 - Select final feature sets to use on test set
- What is the development set for?
 - Parameter tuning (instead of CV)
 - Accuracy estimation (pseudo-test set)
 - More data for training

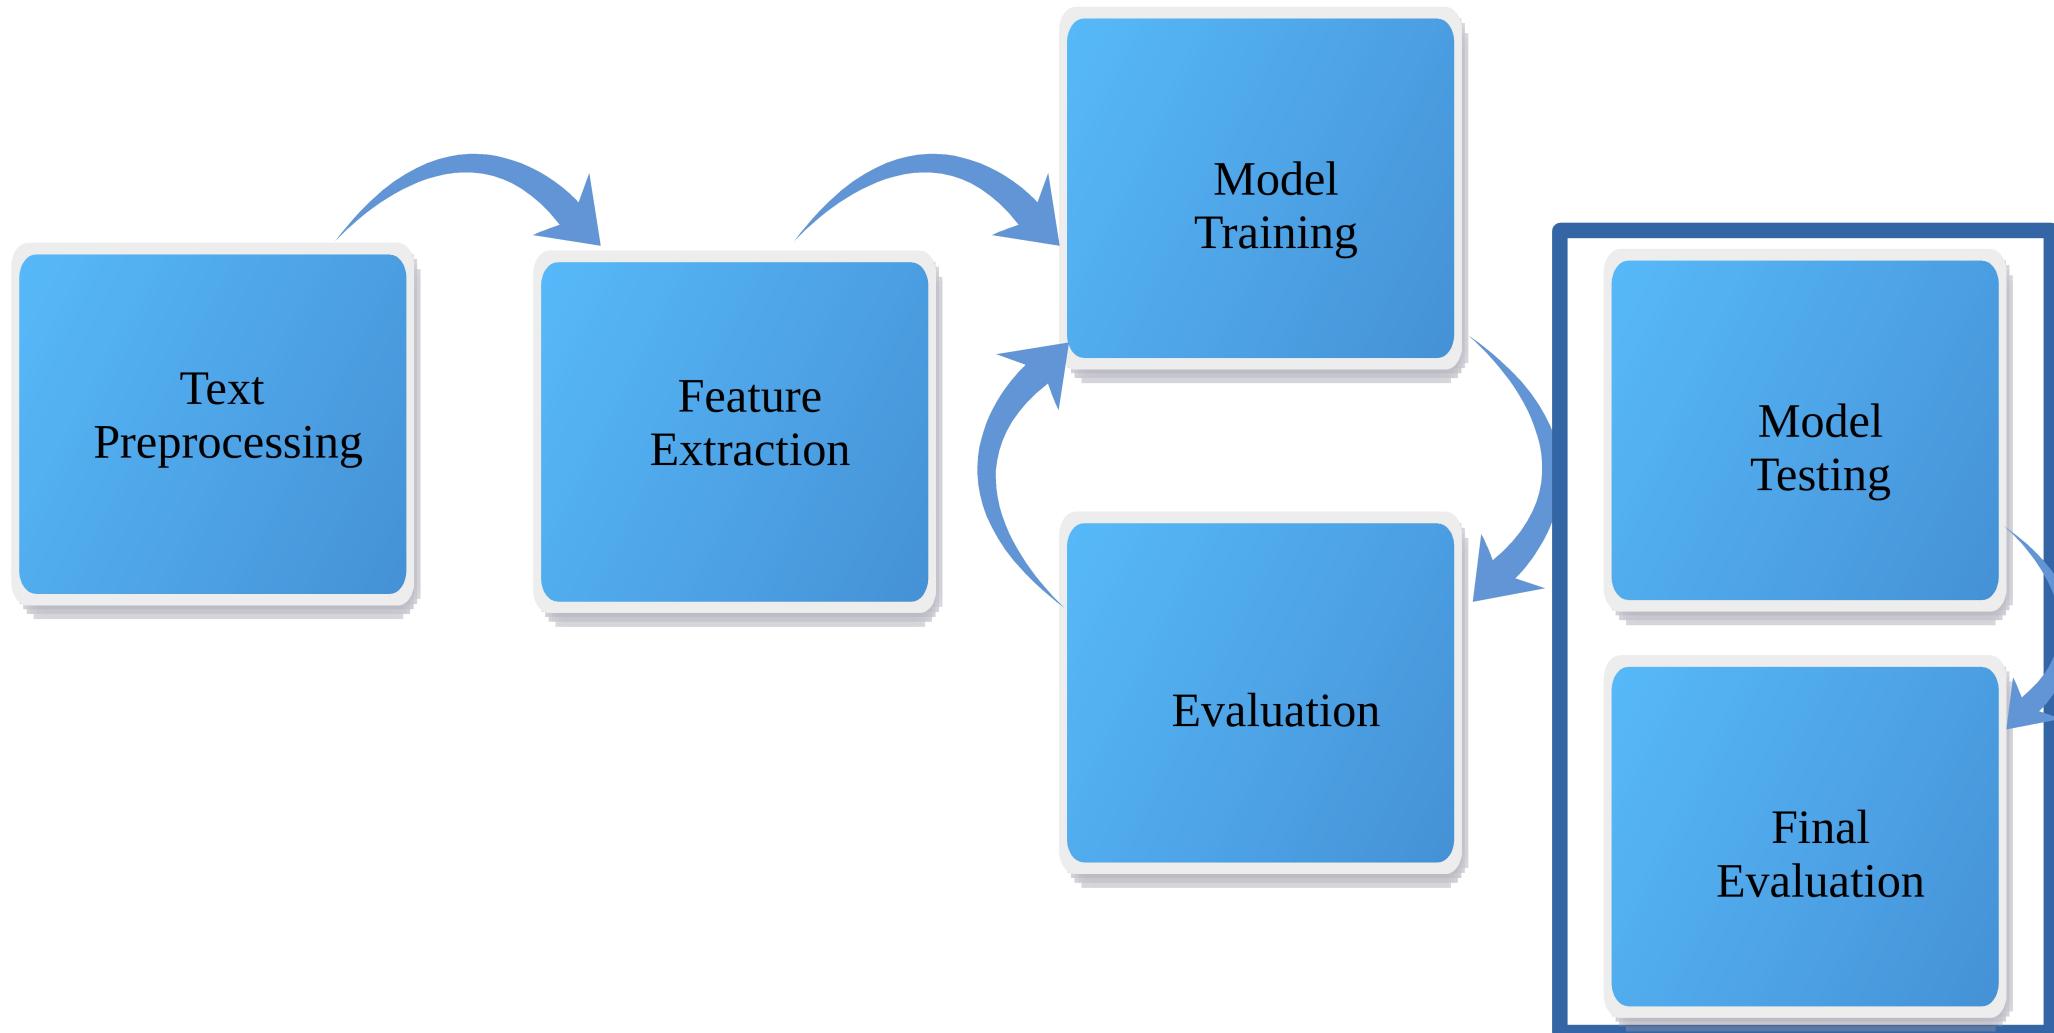
Step 3: tuning dangers

- **Overfitting**
 - You do NOT know the class distribution of the test sets
 - It might be very different compared to your training data
 - Do NOT try to overfit your models
 - a small accuracy boost in 10-fold CV/dev set might result into a large decrease in the test set!
 - **Evaluation:** macro-avg f-score of positive/negative classes
 - Do NOT ignore the neutral class!
 - a) it affects the macro-avg pos/neg performance anyway
 - b) what if the test set has much more/less neutral tweets...?

Step 4: evaluation

- Selected best features/combinations of features?
 - Performance consistent across folds/sets?
 - Different classifiers performing better at different classes?
 - Ensemble...?
 - Error analysis
 - Common characteristics of missclassified tweets?
 - Confusion matrix: frequently missclassified labels?
 - Select best model/parameters/features for the test set
 - Retrain model on training (+dev?) set; only one shot

Exercise 2



Step 5: testing

- Exactly the same pre-processing steps
- Exactly the same features as in the training set
 - do NOT add ngrams you have not seen in the training set
- Simply apply the best model to the test sets
- Make sure you include the results in your submission
- Include detailed evaluation on the three test sets
- Do NOT try to overfit!
 - We will be running your code on different test sets!

Summary

- **Text Preprocessing**
 - Dealing with: user mentions, hashtags, URLs, non-alphanumeric chars, elongated words, all upper-case words, emoticons...
- **Feature Extraction**
 - ngrams, lexicons, word embeddings, twitter-specific
- **Training/Tuning/Evaluation Process**
 - Train different models on different combinations of feature sets
 - **Evaluate!** Error analysis, detailed evaluation, argue about your final choice!
- **Testing Process**
 - Apply model to test sets

Datasets and Marking

- **Training Set:** Extract features, train classifiers, evaluate using 10-fold CV
- **Dev Set:** use for (a) for training, along with Training Set; (b) tuning parameters; or (c) for testing (getting an estimation of performance)
- **Test Set:** Extract the same features & apply your best model
- **Marking** (see the assignment for details!):
 - Make sure your **code** runs (+on different test sets) (20/85)
 - Write a descent **report** (20/85)
 - **Performance** on 5 test sets (20/85)
 - **Feature extraction** approach (15/85)
 - **Methodological** novelty (10/85)

#thank @you :)

Any questions?